



Norwegian University of
Science and Technology

Question analysis of coding questions on StackOverflow

130533 - Knut Lucas Andersen

31-05-2016

Master's Thesis

Master of Science in Applied Computer Science

30 ECTS

Department of Computer Science and Media Technology

Norwegian University of Science and Technology,

Supervisor 1: Assoc. Prof. Simon McCallum

Supervisor 2:

Preface

This is my Master thesis concluding the two years spent at NTNU Gjøvik: Master Applied Computer Science - Web, Mobile, Games track. The thesis was carried out during the spring semester 2016, from January to the end of May.

The main concept for the thesis was based on discussions with supervisor. The original plan was to create a Chat Agent that could answers students questions and give feedback to their question quality, by using StackOverflow as a knowledge base. However, during the Master thesis project presentation, other professors noted that the scope of the project was to large for a Master thesis. The thesis were therefore narrowed down to focus on coding questions posted on StackOverflow, in an attempt to evaluate question quality and predict the future votes for a given question.

31-05-2016

Acknowledgement

I would like to thank the following persons for their help and support during these years. It would not have been possible without them.

My supervisor, Simon McCallum, for understanding my difficult situation and for his patience and helpful advices on how to proceed so that I could complete my Master thesis.

Mariusz Nowostawski for his advice on how to get started with text processing and ideas for the [Support Vector Machines \(SVM\)](#) model.

Rune Hjelsvold for helpful advice in relation to text analysis.

My best friend, Njål Dolonen, for always being there for me, and helped me get through this.

My grandmother, Mimmi H. Underland, may she rest in peace. None of this would have been possible without your support, understanding, love and care. This is for you.

I would also thank my family and friends for believing in me and supporting me through this.

K.L.A

Abstract

Stack Overflow (SO) is today for many developers a well known Question-Answering (QA) system. However, SO has a high requirement to the questions and answers posted, which is reflected through their voting and reputation system. This peer-review processes can be used as an indicator to a questions quality, where questions with high up-votes can be defined as good questions. In this thesis, a system has been developed using Machine Learning (ML) and Support Vector Machines (SVM) to see if it is possible to predict whether or not a new question will be considered as a good or bad question.

This was achieved by using the Stack Exchange (SE) data set, specifically using the one for SO. Questions were divided into two classes, where bad questions was question with a vote score below zero, and good questions were those above zero. Based on content in the various questions, a set of feature detectors was developed and tested against the raw data set. Surprisingly, the features actually lowered the accuracy score (the raw data set had an accuracy of XX%, and the ones using all the feature detectors had an accuracy of XX%).

Add numerical values, and also if time add comparison on those questions which only contained the actual feature

Contents

Preface	i
Acknowledgement	ii
Abstract	iii
Contents	iv
List of Figures	vi
List of Tables	vii
1 Introduction	1
1.1 Problem description	1
1.2 Research questions	2
1.3 Methodology to be used	2
1.4 Justification, Motivation and Benefits	2
1.5 Limitations	3
1.6 Thesis contribution	3
1.7 Thesis structure	3
2 Related work	4
2.1 Stack Overflow	4
2.1.1 Stack Overflows design	4
2.1.2 Stack Overflow and Gamification	5
2.1.3 Stack Overflow and reputation	6
2.2 Asking questions	8
2.2.1 What is the definition of a question?	8
2.2.2 Question classification	8
2.2.3 Text classification	10
2.2.4 Question-Answering	10
2.3 SVM	12
2.4 Dataset	13
2.5 undefined	14
3 Methodology	15
3.1 Dataset and MySQL Database	15
3.1.1 Dataset	15
3.1.2 MySQL Database	15
3.2 Development process	17
3.3 Feature sets, attributes and processing	20
4 Discussions	24
4.1 Data and Testing	26
4.2 Artificial Intelligence Methods	26

4.3	Implementation Architecture	26
4.4	Limitations and other issues	26
5	Conclusion/Summary	27
5.1	Overview of main results	27
5.2	Further work	27
	Bibliography	28
A	Appendix	34
A.1	Acronyms	34
A.2	Data sets/Statistical Overview	34
A.3	MySQL Database	35

List of Figures

1	List of questions, where one can see those with an accepted answers are marked with a green background.	5
2	Example of a question on Stack Overflow	5
3	MySQL Workbench: Setting timeout values to avoid connection loss	16
4	MySQL Database used for dataset	35

List of Tables

1	Results from <code>pandas.DataFrame</code> and <code>pandas.Categorical</code> . -1 is for bad questions (votes < -5), and 1 are for good questions (votes > 50).	21
2	Feature reduction steps before and after text was processed. . . .	22

1 Introduction

Today, many use the Internet as a resource to find answers to their questions and problems. In the past, one was often restricted to only use keywords and not being able to pose the problem as you would when asking another human being. Most search engines today can handle natural language queries, which makes it easier to find the answer you are looking for. The Internet offers a wide range of resources to acquire new knowledge, everything from encyclopaedias to blogs, forums and [Question-Answering \(QA\)](#) communities. One well known [QA](#) community is the [Stack Exchange \(SE\)](#) community, which is built upon the same model as [Stack Overflow \(SO\)](#) [1]. [SE](#) has grown large since its release in 2009, and now contains 154 different communities.

As a developer, one often finds oneself in the situation that a part of the code does not work, you get weird error messages, or you are simply stuck. This is where [SO](#) comes in. [SO](#) is a part of the [SE](#) community, although [SO](#) was actually released before [SE](#). Jeff Atwood and Joel Spolsky wanted to offer programmers a [QA](#) site where they could get the answer they wanted without having to read through a lot of text, see others posting "I also have the same issue" or having to subscribe and pay to see the solution [40]. Question (and answer) quality is maintained through the use of a peer-reviewed gamification system, where users are awarded with votes, reputation and badges for their participation [29, 25, 40, 52]. One of the requirements is that the questions should be of good quality [43, 48, 47]. If a question is bad, users can vote to close or delete it (in which the question will be put on hold). A question can be put on hold or closed if they meet any of the following criterias: Exact duplicate (same question has been asked before), off-topic (not related to [SO](#)), unclear what is being asked, too broad (e.g. could write a book about question being asked) or primarily opinion-based [8, 45].

1.1 Problem description

Most of the systems that have been developed so far focuses on finding the best answer to a question asked by the user. Few, if any, focus on the quality of the question being asked. What defines a good question, and can we in anyway predict whether or not a new question posted on [SO](#) will be considered good or bad by the community? There are many users who have either a negative view or relationship in regards to [SO](#). Many experience that their questions get down-voted, closed or even deleted. For some, they simply do not know how to ask an acceptable question. Questions related to homework are one example of questions that

add
ex-
am-
ples
here

are not accepted on [SO](#). There is even a post on Meta.StackExchange discussing whether or not it should be acceptable to use greetings and sentiments in posts [7]. Therefore, the question becomes: What is and is not a valid question on [SO](#)?

1.2 Research questions

- What defines a good (coding) question on [SO](#)?
- Can we predict a questions quality by using [Support Vector Machines \(SVM\)](#)?
- What type of features increases the accuracy of the [SVM](#)?

1.3 Methodology to be used

The theoretical background in this thesis is mainly focused on [Question Classification \(QC\)](#) and similar research in relation to [SO](#). What has been the focus of other researchers, and in what way did they proceed to solve their questions? The analysis of the questions are done by using the publicly available database dump, which is available via [SE archive](#)¹ [41]. There are several others who have used the same dataset [1, 2, 13, 29, 32, 37, 49, 57]. Taking into consideration that [SO](#) was released in 2008, it means that it now contains approximately 8 years of peer-reviewed data. Because of the size of the data set, and the total amount of posted questions, going through all questions manually would be too time-consuming. Therefore only a select few were studied too see if it was possible to identify what separated the highly up and down-voted questions.

The goal was to develop a [Machine Learning \(ML\)](#) learning system which was based on [SVM](#), since many papers document that this has the best classification accuracy for text classification. The methodology therefore also includes a documentation on the development process, and how and why the given features used were selected.

For the sake of replicability, and also be able to undo potential errors, the system is available in a a GitHub repository². In addition to the source code, the repository also contains both the samples that was used (stored in CSV files), and the models that was created.

1.4 Justification, Motivation and Benefits

Many systems focuses only on finding a good answer, and does not ask if it is a good question. As a famous Norwegian saying goes³: "A fool may ask more than ten wise men can answer". This means that new research possibilities could be opened up in relation to researching question quality by expanding the system.

¹StackExchange dataset: <https://archive.org/details/stackexchange> (Downloaded 30. March 2016).

² GitHub repository: https://github.com/klAndersen/IMT4904_MasterThesis_Code

³ Although its origin comes from a Danish word collection from 1682: https://snl.no/En_d%C3%A5re_kan_sp%C3%B8rre_mer_enn_ti_vise_kan_svare.

Since all the communities within [SE](#) is based on the same model, few modifications would be needed to scale the program to be used within the other communities. As noted in several papers [[29](#), [30](#), [32](#), [25](#), [52](#), [57](#)], question quality is measured based on the amount of votes given. Which can also be compared against the peer-review process in academia, and given that [SO](#) targets professionals and experts, using [SO](#) as a scientific reference is not that far-fetched⁴. [SE](#) has also been the focus of various researchers these past years [[53](#)]. Improving ones own ability to ask better questions can also have a pedagogical effect, which means that this system could be implemented in education.

1.5 Limitations

The greatest limitation is the time available. A large amount of time was spent on setting up the database, and retrieving the questions (the Posts table contains both questions and answers). Only a selection of the questions were selected (a total of 20,000 questions), and training the [SVM](#) over one sample set can easily take several hours. This also has an impact on classification accuracy, since in some cases there is only a small amount of the questions that contains a given feature (e.g. the hexadecimal feature, which only was present in 160 of 20,000 questions). A limitation is also that the focus is only on [SO](#), which means that one would need to make additional adjustments and add more filtering to account for the differences that may occur in each community.

1.6 Thesis contribution

This thesis contribution can be summarized as to the following: Predicting (programming) question quality by using [Artificial Intelligence \(AI\)](#) and [ML](#) to improve the questions quality. Instead of posting bad questions that can get downvoted or closed, the developed system could be able to give feedback to the questions quality. Furthermore, the research presented could open up for new research in relation to how we ask questions online, and in what ways these best can be analysed. It can also be used for educational purposes, e.g. having questions iteratively improve their question quality by asking the system questions.

1.7 Thesis structure

The following is the structure of this thesis:

- Chapter [2](#): State of the art and relevant research
- Chapter [3](#): Methodology
- Chapter [4](#): Discussion on development, the thesis and limitations
- Chapter [5](#): Conclusion and suggestions for further work

⁴ Posnett et al. [[32](#), p. 1] noted that [SO](#) "ranked 2nd among reference sites, 4th among computer science sites, and 97th overall among all websites".

2 Related work

2.1 Stack Overflow (SO)

2.1.1 Stack Overflows (SO) Design

The following list shows the design used in SO (based on M. Sewak [25, p. 6-7] and Treude et al. [52, p. 805]):

1. Votes: Questions and answers which are considered good (or bad) by the community can be given a score. This gives a filtering mechanisms, which allows users to ignore answers that are bad or wrong. Furthermore, answers are sorted by votes, and you can also sort questions on SO by vote score (see Figure 1).
2. Accepted answer: If the user asking a question gets an answer that they find satisfactory, they can select it as the "accepted answer". This answer will be the first displayed of the answers, and is also viewable when searching for questions (see Figure 1 and 2).
3. Tags: Each question is associated with a tag¹, which can be a topic, a programming language, a methodology, etc.
4. Badges: Similar to achievements in games, Badges are used to reward the user for their participation.
5. Reputation and Bounty: Currency system for user participation. E.g. voting for questions, getting your answer selected as the accepted one, etc. Bounty is a trade, where if your question goes unanswered for too long, you offer up parts of your reputation to receive an answer.
6. Data dump: Available data dump containing all content available within the SE community [41]. You can either download single files, or everything by using a Torrent client.
7. Pre-Search: Encouraging users to check that their question is not already posted by presenting a search bar when asking a new question.
8. URL keywords and Google: The questions title is included in the URL, allowing it to be processed by search engines. In addition, Google uses their crawlers every 10 second to have the latest updates from in their search engine [12].
9. Critical mass: Before Atwood and Spolsky launched SO, they invited developers and programmers to participate to have some domain experts available.

¹ A full list can be seen here: <http://stackoverflow.com/tags/>




Figure 1: List of questions, where one can see those with an accepted answers are marked with a green background.

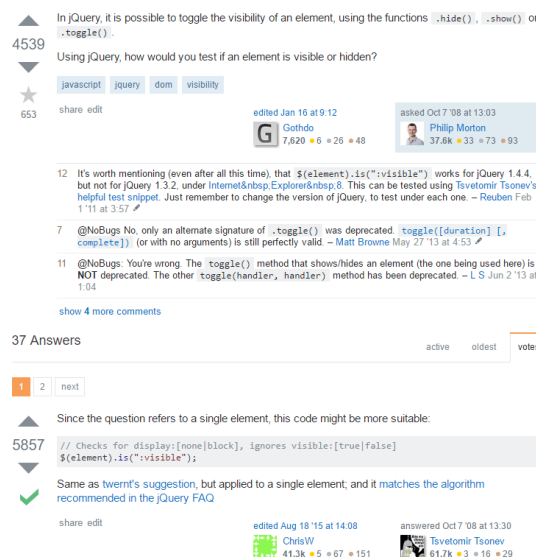


Figure 2: Example of a question on Stack Overflow²

2.1.2 Stack Overflow (SO) and Gamification

Deterding et al. [9] defines Gamification as "the use of game design elements in non-game contexts", and is the definition this section will be based on. Several papers make notes of the pedagogical and educational aspect of SO [30, 32, 57], and [30, 57] use the term gamification in their paper. One of the founders, Jeff Atwood said in an interview that he wanted users to not just give good answers, but also trick them into improving their communication skills [32]³. In

²Source: <http://stackoverflow.com/questions/178325/checking-if-an-element-is-hidden>

³ From this interview:
<http://www.wired.com/2012/07/stackoverflow-jeff-atwood/2012>.

the course IMT4007 Serious Games Simon McCallum and Marius Nowostawski, presented their game GoRad, which was based on us students reading articles and posting questions which were voted on. The SO system awards users based their activity by using votes, reputation and badges [25, 29, 52, 42, 46]. In relation to gaming, there are four player types: Achievers, Explorers, Socializers and Killers [26, p. 3].

These player types can be used as a representation⁴ for the various users of SO. Achievers are there for the reputation and badges, socializers are to interact, discuss and share knowledge. Explorers might find joy in looking at various topics, or searching for unanswered questions. The only exception would be the "Killer" type. Killers are those "... who always want to create trouble/problems for other participants" [26, p. 3]. In an online QA system (or Internet in general), these are what are commonly referred to as "Trolls" [11, 4]. However, due to the system used in SO, Trolls would not be able to survive, simply because the reputation controls what you have access to [44]. If you down-vote a post, you lose reputation. If your post gets down-voted, you also lose reputation. Users who are not willing to follow the guidelines can be locked out of SO [3]. However, today there is a lot of blogs complaining about the current structure of SO, who claims that a lot of the moderators are trolls⁵.

2.1.3 Stack Overflow (SO) and reputation

Many QA sites includes domain experts to ensure some quality is upheld, and uses voting and reputation as a quality measurement [2]. Furthermore, questions topics, page views and votes can be used by search engines as a ranking mechanism, and it helps users to find the answers they are looking for. Anderson et al. [2] identifies two principles for the answer process. This process starts with the question being filtered down through the users, starting with domain experts. If the domain experts does not answer, it goes further down the chain, until it in the end either gets an answer, or is not answered at all. Both Anderson et al. [2] and Treude et al. [52] defines an unanswered question to be a question where no accepted answer is chosen⁶. The second principle is that a questions activity level does not just indicate the interest for the question, but could also be an indicator for quality (because a question can have multiple answers).

⁴ Yang et al. [57] characterised users as "Sparrows" and "Owls", where sparrows answers question for reputation and owls answers the difficult ones (domain experts).

Ahmed et al. [1, p. 2] defined users as "lurkers, help-seekers (askers) and givers (responders)".

⁵https://www.reddit.com/r/programming/comments/3cafkp/is_stack_overflow_ouerrun_by_trolls/.

<https://medium.com/@johnslegers/the-decline-of-stack-overflow-7cb69faa575d>

Last accessed 23.05.2016.

⁶ However, they do not take into considerations users who find a solution on their own, or simply forget or neglect to mark a an answer as accepted.

Since users can only gain 200 reputation points daily, the only way to earn more is by having your answer marked as accepted or through bounties [46]. Movshovitz-Attias et al. [29] found that users earn more reputation by providing good answers rather than good questions⁷. Most questions were asked by the users with a low reputation, but on average users with high reputation asked more questions. This indicates that reputation could be used as a measurement for expertise. Ahmed et al. [1] also found that there was a correlation between amount of answers given and the users reputation.

Yang et al. [57] found that the activity level of a user is not equal to knowledge, and divided users into two groups; "Sparrows" and "Owls". The sparrows are the basic users who earn reputation and badges by answering the easy questions, and has a greater interest in the gamification element. They found that the sparrows usually have a low average score and target questions that are easy, or non-relevant. Nonetheless, they are still important since they are able to provide quick feedback. As for the owls, they are considered to be the domain experts. The owls earn reputation by asking more advanced questions, providing better answers (i.e. getting their answer accepted) and answering popular and difficult⁸ questions.

Posnett et al. [32] views SE and SO as a learning community, since users help each gain new knowledge, and motivates learning. They wanted to see if the quality of the users answers improved over time. By constructing a posting history for each user, they found that the overall answer score decreased, and that the answer quality was static.

Nasehi et al. [30] did a qualitative analysis of code examples posted on SO. Their focus was on questions related to Java programming, with the requirements that the question should at least have a score of +4 and the answer +7. In addition, a code example should be included (by checking for `<code>` in the post). They found that the code explanation was just as important as the code examples (but you are still restricted to the quality of that example). For the code to be considered good, they listed the following attributes:

1. Concise code: Code samples should not be too long. They should be simple, and only focus on the parts that are relevant to the topic. Additional or non-relevant parts should instead be documented by using descriptive comments.

⁷ However, as stated in Movshovitz-Attias et al. [29, p. 3], the reputation system was changed at one point. Originally, up-votes on questions and answers gave users a +10, but this was later changed into up-votes on questions only giving +5.

⁸ Popularity was measured based on page views and the time between a question was posted until an answer was selected as accepted. The popularity can also therefore be seen as a measurement for difficulty. The longer it takes to answer, the more difficult the question is. [57, p. 273].

2. Question context: If the code is not working properly, suggestions for improvement should be added. One could also explain best practices and suggestions for improved readability. This will also have a pedagogical benefit, since the user asking the questions will learn to write better code.
3. Highlighting important elements: "Straight to the point", clearing up misunderstandings, pointing to relevant resources, etc.
4. Step-by-step solution: Splitting code into chunks, and explaining each chunk and its functionality. Comparison of languages; e.g. "How can I do X in C#, when I'm used to Java?"
5. Providing links to extra resources: Answers can be kept short by adding links to external resources, but a short summary should still be added.

2.2 Asking questions

2.2.1 What is the definition of a question?

The context of a question varies within the setting it is used. A question can be broad, where multiple answers can all be correct, or they can be factual, having only one right answer. When you are asking someone a question, you ask because you want to either find a solution to a problem, or learn something new. In the context of learning, questions are used for evaluating the students knowledge, or help them learn something new Nielsen et al. [31].

When doing research, you need research questions and hypotheses to decide what the goal of your research is. What questions are you trying to find an answer to, and what does that answer tell you? Slowiaczek et al. [38] defines asking a question as information selection and the answer(s) to a question as information usage. If you are working with statistical data, and you just post the numbers, this will not inform anyone. You need to explain what the numbers mean, and how you got them. The quality of an answer is also restricted to the quality of the question you ask. One can therefore assume that if you ask a good question, you will get a good answer [38].

2.2.2 Question Classification (QC)

QC is the process of categorizing a question into a class or category based on its structure, usually to decide what the expected answer type is [22, 23, 24]. To classify a question, it is important to select only those features that helps you identify the class it belongs to. To get a classification results, you use what is known as a classifier. The quality of a classifier can be measured by its accuracy and precision (see Equation 2.1 and 2.1; taken from [22, p. 13]).

$$\text{Accuracy} = \frac{\# \text{ of correct predictions}}{\# \text{ of predictions}} \quad (2.1)$$

$$\text{Precision}[c] = \frac{\# \text{ of correct predictions of class } c}{\# \text{ of prediction of class } c} \quad (2.2)$$

WH-words

WH-words are mostly found in factoid questions [24]. Huang et al. [15] listed eight different WH-words: What, which, when, where, who, how, why, and rest (rest being the type does not belong to any of the previous type). Letovsky [20] also listed "Whether" and "Discrepancy"⁹. However, not all are equally easy to use for classification, because even if the questions ask for the same answer, wording and syntactic structures can make it difficult to classify. Question containing words like "What", "Why", "How" and "Which", can be harder to classify due to the lack of limitation in regards to answer types¹⁰ [15, 24].

Bag of Words (BOW) and N-grams

N-gram is a model that is used for splitting text into either characters (character model) or word frequencies (word model). The **Bag of Words (BOW)** model (or unigram) only looks at singular words, ignoring the order and relies only on the frequency for each word [28, 33]. Bi-grams takes dual values, tri-gram takes three, etc.

One problem with N-grams is that the dimension of the feature space is equal to the amount of words in the vocabulary [33, 23]. When using categorization, there can be issues with mapping new words that does not exist in the vocabulary [58]. The impact of N-gram is also related to the size of the text being analysed. Zhang and Lee [59] found that there was not a big difference when using between bag-of-ngrams (all continuous word sequences in the question) and BOW as features.

Word mapping and processing: Case-sensitivity, Stemming, Stop words and Tokenization

To reduce the amount of words used, there are more steps that can be taken. By removing the case-sensitivity, all words will be equal (e.g. is the word 'Hello' equal to the word 'hello?'). [15] includes case-sensitivity under a definition called word shape, consisting of five elements: upper case, all lower case, mixed case, all digits, and other.

Semantics can be used for word filtering, e.g. removal of duplicate words or words with same meaning. WordNet has a built in function called synsets() which removes synonyms (words having the same meaning). You can also look for hypernyms (words belonging to a category with a parent-child relationship) or use stemming. Stemming reduces the word to its base-form, e.g. crying would be converted into the word cry. Word separation is also possible through tokenization, which splits the text into an array based on a set delimiter. There is also usage of stop words for removal of frequently used words in a given lan-

⁹ "Questions that reflect confusion over a perceived inconsistency." [20, p. 5]

¹⁰ An answer type (or named entity) is the expected type of the answer to a given question (e.g. a Location, Organization, Person, Date, etc) [14, 24, 34, 58].

guage.

Grammatical properties can be extracted by using [Part of Speech \(POS\)](#), e.g. by using [Natural Language Toolkit \(NLTK\)](#)¹¹, which can be helpful in reducing ambiguities [5]. Li and Roth [22] uses the word head chunks to identify what the question is asking for when multiple types are introduced (avoid ambiguity). The same concept is used in [15] and [23], but there it is referred to as headwords.

2.2.3 Text classification

The goal of text classification (or text categorization) is to be able to process multiple documents or large amounts of text into categories. It shares similarities with question classification, although an obvious difference would be the size of the text that is processed. Some examples are spam filtering [33, 18], to identify languages, or filing documents based on content [18]. Documents can belong in more than category, and since categories can overlap they must be treated as a binary classification problem [17]. Text classification starts with retrieval of the documents, usually by using [Information Retrieval \(IR\)](#) methods, and then transforming the text into features for the classification. When you have a large amount of text, you can easily get a feature space that is very high dimensional, and that is why feature selection is important. Feature selection is the selection of features (or attributes) that are important for the classification. E.g. if you were classifying documents based on colour description, then hypernyms for colour would be an important feature.

2.2.4 Question-Answering (QA)

[QA](#) is mostly used as a method for finding the answer to a question from an unknown amount of documents. When using a search engine, one can accept that there are several results that are listed because at least one of the search terms exists. However, when using a [QA](#) system, users want the answer straight away instead of having to read through several documents. In addition, [QA](#) sites allow users to search for questions in the same way they would ask another human (natural language¹²), and there are also different types of [QA](#) sites. Domain specific [QA](#) focuses on a specific topic (e.g. [SO](#)) and open domain where everything goes. [QA](#) sites can also function as an archive or a knowledge base, since all the posts are available even years after they were posted.

Yen et al. [58] found that it was more efficient searching for answers in a small dataset, then the document as a whole. By using a passage retriever, the documents were split into paragraphs and ranked by using evaluation metrics. Isozaki [16] could not use TF-IDF based paragraph retrieval, because the paragraphs

¹¹ [NLTK](#) includes in their [POS](#) tagger the following grammatical properties: Adjective, adposition, adverb, conjunction, determiner, article, noun, numeral, particle, pronoun, verb, punctuation mark and others [50, See Section 2.3].

¹² However, there is a problem with linguistics in natural language systems [24].

were too short to cover all query terms. If the terms that were used were too short or too long, the passage scores would not reflect the density distribution. Xu et al. [56] built an online QA system for tourism, which consisted of question analysis, information retrieval and answer extraction. Since rule-based approach requires expert knowledge, creating features that are domain specific can improve accuracy (what they called "domain term concept hierarchy"). To validate the classification, they tested the results by using 5-fold cross-validation.

Li and Roth [22] used semantics to categorize questions based on the possible semantic answer type. One issue with questions are that since they can be very short, they contain little text. However, the lack of long text improves both the accuracy and analysis.

Zhang and Lee [59] says that QC is important, and simply looking for WH-words is not enough. By using WH-words, headwords, WordNet semantics, N-grams and word shapes as features, and a linear SVM and Maximum Entropy model, they reached an 89.2% and 89.0% over a standard benchmark dataset. They also experimented with four other algorithms, Nearest Neighbours (simplified version of k-NN), Naive Bayes and Decision Tree and Sparse Network of Windows (SNoW). However, these were outperformed by the SVM.

2.3 Support Vector Machines (SVM)

SVM is the chosen ML algorithm that was chosen to use for the question analysis. This section is mainly intended as a light introduction to what SVM is, and what the most common uses are. Furthermore, in what way can this be utilized for text and question classification?

In QC, question is represented using vector space model, the question is a vector which is described by the words inside it. [23]

SVM is based on the Structural Risk Minimization principle from computational learning theory. The idea of structural risk minimization is to find a hypothesis h for which we can guarantee the lowest true error. The true error of h is the probability that h will make an error on an unseen and randomly selected test example. An upper bound can be used to connect the true error of a hypothesis h with the error of h on the training set and the complexity of H (measured by VC-Dimension), the hypothesis space containing h . SVMs find the hypothesis h which (approximately) minimizes this bound on the true error by effectively and efficiently controlling the VC-Dimension of H . One remarkable property of SVMs is that their ability to learn can be independent of the dimensionality of the feature space. SVMs measure the complexity of hypotheses based on the margin with which they separate the data, not the number of features. This means that we can generalize even in the presence of very many features, if our data is separable with a wide margin using functions from the hypothesis space. The same margin argument also suggest a heuristic for selecting good parameter settings

for the learner (like the kernel width in all RBF network). The best parameter setting is the one which produces the hypothesis with the lowest VC-Dimension. This allows fully automatic parameter tuning without expensive cross-validation. Theoretical analysis concludes that [SVMs](#) acknowledge the particular properties of text:

- high dimensional feature spaces,
- few irrelevant features (dense concept vector),
- sparse instance vectors.

Experimental results show that [SVMs](#) consistently achieve good performance on text categorization tasks, outperforming existing methods substantially and significantly. With their ability to generalize well in high dimensional feature spaces, [SVMs](#) eliminate the need for feature selection, making the application of text categorization considerably easier. Another advantage of [SVMs](#) over the conventional methods is their robustness. [SVMs](#) show good performance in all experiments, avoiding catastrophic failure, as observed with the conventional methods on some tasks. Furthermore, [SVMs](#) do not require any parameter tuning, since they can find good parameter settings automatically. All this makes [SVMs](#) a very promising and easy-to-use method for learning text classifiers from examples. [\[17\]](#)

2.4 Dataset

short about the datasets, others who has used it, and datasets in general, e.g. [\[19, 39, 55\]](#)

2.5 undefined

Stanley and Byrne [\[49\]](#) - Predicting Tags for StackOverflow Posts

Short et al. [\[37\]](#) - Tag Recommendations in StackOverflow

Wang et al. [\[54\]](#) - An Empirical Study on Developer Interactions in StackOverflow

Lezina and Kuznetsov [\[21\]](#) - Predict Closed Questions on StackOverflow

3 Methodology

3.1 Dataset and MySQL Database

3.1.1 Dataset

The dataset contains all information that is currently available in the [SE](#) community (at the time the dataset was created). The following is a list of the tables found in the dataset:

- Badges: Badges awarded to users.
- Comments: Comments given either to a question or an answer.
- Posts: Posts on [SE](#), this contains both questions and answers.
- Posthistory: The history of a given post (e.g. edits, reason for closing, etc.).
- Postlinks: Link to other Posts (e.g. duplicates).
- Users: Information about the given user registered at the given community.
- Votes: Type of vote given to a Post (e.g. up/down, vote to close, etc.).

In the beginning, the dataset that was used was downloaded in August 2015. However, since this turned out to be outdated, the latest dataset was downloaded from (<https://archive.org/details/stackexchange>) on 30. March 2016. The dataset comes in zip-files, where each zip-file contains all the rows found in the given table. These rows are presented in an XML file, as shown in Listing 3.1.

Listing 3.1: Content in stackoverflow.com-Tags.xml

```
<?xml version="1.0" encoding="utf-8"?>
<tags>
<row Id="1" TagName=".net" Count="227675"
ExcerptPostId="3624959" WikiPostId="3607476" />
<row Id="2" TagName="html" Count="511091"
ExcerptPostId="3673183" WikiPostId="3673182" />
...
</tags>
```

3.1.2 MySQL Database

In the beginning, the issue was getting access to the file and see how it looked like. Since most of these XML files had a large file size (ranging from 3,9 MB to 71,9 GB) none of the editors could open them. Attempting to open them through Python code also failed, since there was not enough memory to process everything. The only solution was therefore to create a MySQL database that could contain all the data.

Setting up the MySQL database was not a straight forward process. The operative system I was running was Arch Linux, where they had switched from using Oracle's MySQL to MariaDB¹. One of the main problems was the available storage space² and the varying file sizes. Some of the issues were mainly connection timeout, no more disk space and connection loss (e.g. "Error Code: 2013. Lost connection to MySQL server during query"). To avoid losing the connection to the database, the timeout values had to be changed in MySQL Workbench (shown in Figure 3).

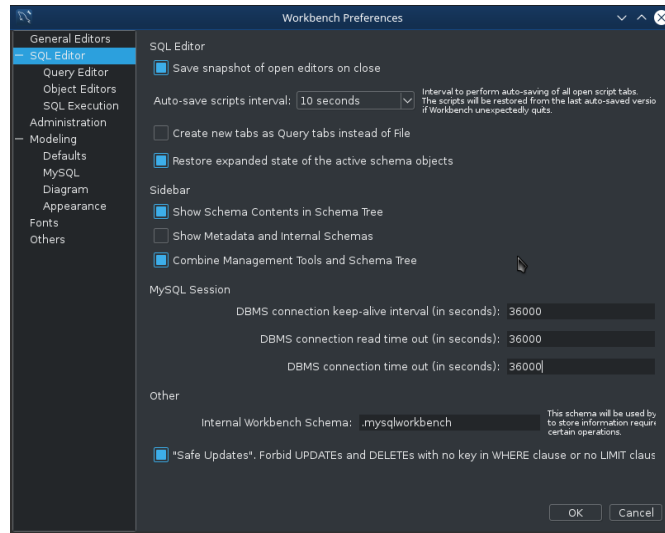


Figure 3: MySQL Workbench: Setting timeout values to avoid connection loss

The next problem was the lack of disk space. MySQL by default stores all databases and belonging tables in `/var/lib/mysql/`, and it also creates temporary backup files (where the file size is equal to the size of the current database). Since the default folder for temporary files was on `/root`, the disk space was used up in less than 30 minutes. Therefore, two things needed to be done. First, disable the storage of temporary files, and secondly change the storage location for the database. The problem when tinkering with the configuration file is that things easily break. Which is what happened, and a clean install was needed for both MariaDB and MySQL (the changed settings can be seen in Listing 3.2). The final step was to create symbolic links that linked the database to the location where the tables were stored (this has to be done before creating the tables, if not MySQL Workbench will store the tables in `/var/lib/mysql/`)³.

¹ See <https://wiki.archlinux.org/index.php/MySQL>.

² The HDD with Arch Linux installed had a disk size of 500 GB, with four partitions; root, var, swap and home. 40 GB was used for `/root` and `/var`, 12 GB was used for swap and the remainder was used for `/home`.

³ It should be noted that after an upgrade of MariaDB, MySQL and MariaDB could no longer

Listing 3.2: Changes made to config file: /etc/mysql/my.cnf

```
# disable storage of temporary files
#tmpdir = /tmp/
# disable storage of log files
#log-bin = mysql-bin

# set directory for storing database files
datadir = /home/mysql
```

Listing 3.3: Load XML file into a table in the MySQL database

```
LOAD XML LOCAL INFILE
path_to_xml_file
INTO TABLE db_table
ROWS IDENTIFIED BY '<row>';
```

Listing 3.3 shows how the files were loaded into the tables, and the complete database can be seen in Appendix A.3, p. 35. Since the Posts table is large (~29,5 million rows) and it contains both questions and answers, two new tables were created; "posvote_Posts"⁴ and "negvote_Posts". posvote_Posts contains questions with a score higher then zero (score > 0) and negvote_Posts contains all questions with a score lower then zero (score < 0).

3.2 Development process

When starting the development, the focus was on retrieving the data from the database, and processing it for text analysis. To be able to store all the retrieved columns and the belonging rows without creating object classes, the pandas.DataFrame⁵ was used.

The questions retrieved needed to be processed before any analysis could be done. The reason for this is because the questions was written as HTML (including HTML entities). An example is shown in Listing 3.4. Every question starts with the <p> tag, and if the question contains code samples, these are wrapped with a <code> tag. To convert the HTML text into readable text, a HTML parser class was created (based on answer by [10]).

find the tables, even if they still were in the /home/mysql/ folder. It is therefore advisable to dump the database after inserting all the tables, since it goes a lot faster to restore the database from dump rather than insertion from XML files.

⁴ The Posts table has a file size of ~43,6 GB, whereas posvote_Posts file size is ~11,2 GB. negvote_Posts has a file size of ~1,33 GB.

⁵ Pandas: <http://pandas.pydata.org/>.

Listing 3.4: Question before HTML is removed (Question ID: 941156)

```

<p>
Why do we need callbacks in ASP.NET or any server side technology?
</p>&#xA;&#xA;<p>One answer can be, to achieve asynchronous calls.
</p>&#xA;&#xA;<p>But I am not satisfied with this answer.</p>
&#xA;&#xA;<p>Please explain it in a different way.
</p>
&#xA;

```

To process the questions, CountVectorizer from scikit-learn was used. CountVectorizer uses the vocabulary found in the text and counts the frequency of each word [36] [35, see 4.2.3]. When looking at this vocabulary, a lot of unimportant words was found (a lot which came from the code samples) in some of the questions. At first all code samples were removed from the text, but later on they were replaced with the value 'has_codeblock', indicating that this question contained one or more code samples. This was achieved by using a combination of lxml⁶ and bs4⁷ (BeautifulSoup). lxml was used to construct an XML tree containing all the tags (to be able to retrieve the content by searching for a given tag), and bs4 was used for beautifying the HTML (since in some cases an error was thrown complaining about "Missing end tag").

However, for some questions, part of the text was lost, and for others, some <code> tags was not removed. On inspection, it was found that the trailing text following the <code> samples was stored in a .tail attribute. Since the <code> was removed, the .tail attribute was also removed. This was fixed by storing the content of the <code> .tail attribute into its <parent>⁸ (where <parent> is the tag that contained the given <code></code>) .tail attribute. As for the non-complete removal of <code> tags, this error mostly occurred for code samples that contained XML or HTML code⁹, because the lxml parser failed. The solution was to replace the lxml parser with bs4 and just change the content of the <code> tag to the value 'has_codeblock'.

Considering the size of the dataset, and that the source code was hosted on GitHub, I was hesitant to store the training data in a separate file. However, when loading 20,000 samples from the database with a 'WHERE' parameter, things tend to go more slow. At this point, it was decided to try to dump the loaded data from the database to a file. This was achieved by using pandas.DataFrame.to_csv¹⁰.

⁶lxml: <http://lxml.de/>

⁷BeautifulSoup: <https://www.crummy.com/software/BeautifulSoup/>

⁸ It was also necessary to check if the <parent> had a .tail, if not, the .tail attribute had to be set for the <parent> to avoid the error: "NoneType + str: TypeError".

⁹ One example is this question:

<http://stackoverflow.com/questions/19535331/print-page-specific-area-or-element>.

¹⁰ pandas.DataFrame.to_csv:

At a later point, the unprocessed dataset was also dumped to a CSV file for replicability¹¹.

Further examination showed that the vocabulary contained a lot of numerical and hexadecimal values, but also a lot of non-English words. The numerical and hexadecimal values were replaced using regular expressions to 'has_hexadecimal' and 'has_numeric'. The non-English words were a bit more troublesome to handle, since these were mainly used to prove a point or show an example of the issue they were having¹². Attempts were made to filter them out by using `corpus.words.words()` and `corpus.wordnet.synset()` from NLTK¹³, and PyEnchant¹⁴. However, WordNet does not have a complete database of all English words, and they all claimed some words were not English even though they were. The solution turned out to be a lot simpler. Instead of creating filters, the CountVectorizer already had one built in. By adjusting the minimum document frequency (`min_df`) and setting it to 0.01, words that appeared in less 1% of all documents were ignored.

To be able to run the system without relying on an [Integrated Development Environment \(IDE\)](#), making it run from the Terminal using basic command setup seemed like a good idea. At first `optparse` was used, which ironically turned out to be deprecated and replaced by `argparse`. However, the problem was that you could only run one command at a time, whereas I wanted the program to be able to run until exited. The reason for this was because it needs to load a model before it can make a prediction, in addition the user might want to predict multiple questions. This was therefore replaced with a basic while loop that runs until the user enters the exit command. The setup used for `argparse` was kept, so users from *nix system might be more familiar with similar commands (shown in Listing 3.5). At the end, there were some commands that were not added to the menu, since these were mostly used for testing.

Listing 3.5: System menu

Menu:

d: Loads default model (if exists) from `./pickle_models`

e: Exit the program

h: Displays this help menu

l: Load user created model. Arguments:

 path: Path to directory with model(s) (e.g. `/home/user/my_models/`)

 filename: The models filename

http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.to_csv.html.

¹¹The only change made to the unprocessed dataset was removing the HTML tags.

¹²<http://stackoverflow.com/questions/856307/wordwrap-a-very-long-string>.

¹³<http://www.nltk.org/>

¹⁴<http://pythonhosted.org/pyenchant/>

suffix: File type – Optional (default: '.pkl')

p: Predict the quality of the entered question. Arguments:
 question: Question to predict quality of

t: Train a new model based on an existing (or new) data set. Arguments:
 path: Path to directory with training data (e.g. /home/user/my_data/)
 filename: Filename for data set (model name will be the same as this)
 db_load: Load from database (Enter 0: No, 1: Yes)
 limit: Limit for database row retrieval (integer) – Optional unless 'db_load' is '1'

u: Create an unprocessed data set based on database content (from database set in dbconfig.py).
 Arguments:
 filename: Filename for data set (model name will be the same as this)
 limit: Limit for database row retrieval (integer)
 feature_detectors: Create singular feature detectors based on data set? (Enter 0: No, 1: Yes)
 create_model: Create classifier model(s) based on data set?
 0: No, 1: Unprocessed model, 2: Feature detector model(s) 3: Both (1 and 2)

3.3 Feature sets, attributes and processing

When retrieving the questions from the database, the vote score was set to less than -10 for bad question and greater than 50 for good questions (retrieval limit set to 10,000; 20,000 total). However, the vote score was set too low for the bad questions, since only 683 rows was returned. Therefore, the score was then set to less than -5. What was also found when using `pandas.Categorical` to get an overview (code snippet in Listing 3.7 and result in Table 1), one can see that for the 10,000 bad questions, the average vote score was -7. This could be an indicator that when a question has a vote score below -5, they are ignored.

Class	Statistics	AnswerCount	Score	Question length
-1	mean	2.0483	-7.0275	319.226
	std	1.3129	2.676	382.115
	min	0.0	-147.0	13.0
	25%	1.0	-7.0	153.0
	50%	2.0	-6.0	239.0
	75%	3.0	-6.0	379.0
	max	20.0	-6.0	13673.0
1	mean	11.9379	182.5483	459.329
	std	13.707824	317.47217	531.187559
	min	0.0	51.0	13.0
	25%	6.0	67.0	189.0
	50%	9.0	96.0	328.0
	75%	14.0	173.0	558.0
	max	518.0	9432.0	18867.0

Table 1: Results from pandas.DataFrame and pandas.Categorical. -1 is for bad questions (votes < -5), and 1 are for good questions (votes > 50).

Listing 3.6: Getting Categorical data from pandas.DataFrame

```
from pandas import DataFrame, Categorical

# get statistics from pandas.DataFrame
temp_df = __so_dataframe.loc[:, ("Score", "Body", "Title",
                                "AnswerCount", "length")]
temp_df.loc[:, CLASS_LABEL_KEY] = Categorical(__so_dataframe.loc[:,
                                                                "label"])

# prints out the questions AnswerCount, Score and length
print(temp_df.groupby("label").describe())
# prints all selected columns
print(temp_df.groupby("label").describe(include='all'))
```

To be able to develop some theories on what the difference between good and bad questions was, a total of 200 questions were reviewed (by sorting questions based on votes¹⁵). It was easier to see certain patterns in down-voted questions rather than those that were up-voted. A repetitive pattern was that many had either no code example, or poorly written code. These questions could also show

¹⁵<http://stackoverflow.com/questions?sort=votes>

Step	Text processing	Vocabulary count	CountVectorizer
1	None	69766	analyzer="word"
2	Stop words	69462	analyzer="word", stop_words="english"
3	Removal of code, hexadecimal and numerical values	27624	analyzer="word", stop_words="english"
4	Minimum document frequency	440	analyzer="word", min_df=0.01, stop_words="english"

Table 2: Feature reduction steps before and after text was processed.

indications of not having tried anything, or that they were based on either homework or school assignments. This in turn lead to a hypothesis that if a question contains indicator of word synonyms for homework¹⁶, it would be considered a bad question. In addition, some code examples had syntax errors, which made the minimum working example (MWE) not executable. Some questions also contained links, either to external resources or indicators of potential duplicates. Therefore links was also considered a potentially useful feature. Tags was also considered as a feature, which was divided into two: Attached and External tag. Attached tags are tags which the user has linked to the question, whereas external tags are all the tags available on SO. Version numbering was also considered, but this was not included due to the complexity of writing a proper filtering method to account for all possible variations.

Features were added in the same manner as was done for code samples, numerical and hexadecimal values. However, there were some issues when attempting to replace the tags and the synonyms for homework. At first, WordNet was used for synonyms (using `wordnet.synset()`). The only problem was that for the word 'homework', `wordnet.synset()` only returns ['homework', 'prep', 'preparation']. Whereas Thesaurus¹⁷ had a lot more suggestions, and was therefore used instead. Words were selected based on whether or not it was plausible that they could be used in programming related question setting. A new problem now arose, namely the issue that the word "assignment" did not necessarily need to occur in a homework setting, since it could also be used as a programming word (e.g. assignment operator¹⁸). Therefore features for homework were split into two types: 'has_homework' and 'has_assignment'.

Tags were without a doubt one of the most annoying features to detect and replace. Site tags (or external) are single text values in the database, whereas the question can have up to five tags attached. Those attached tags are then

¹⁶<http://www.thesaurus.com/browse/homework>

¹⁷<http://www.thesaurus.com/browse/homework>

¹⁸<http://stackoverflow.com/questions/5368258/the-copy-constructor-and-assignment-operator>

separated in the following format: "<c><multi-threading>", which had to be processed by removing the '<' and the '>'. After the removal, each tag value was added to a list, so that all attached tags was indexed based on the question they belonged to. Furthermore, a combination of string replacement and regular expression was needed. The regular expression was used for single character tags (e.g. 'C'), and word replacement for longer words. The reason for this was that when using string replacement, single character tags replaced occurrences even if they appeared in the middle of a word. If the tags contained characters that could be interpreted as a regular expression (e.g. C++), it would give error about multiple repetitions. In addition, the tags needed to be sorted based on their length, since for questions that contained tags which included both <C> and <C++>, if <C> came first, it replaced the <C++> with 'has_*_tag'++. The text also had to be converted to lower-case to ensure proper tag matching.

Listing 3.7: Replacing tags in the question

```
for word in word_set:
    if len(word) == 1:
        # if its only one character (e.g. 'C'), ensure that it is a singular word by using regex
        text = re.sub(r"\b%s\b" % word, replacement_text, text, flags=re.IGNORECASE)
    else:
        text = text.replace(word, replacement_text)
```

4 Discussions

For Discussion chapter: _____

In the **BOW** model, only single words or word stems are used as features for representing document content. The issue is that learning algorithms are restricted to detecting patterns in the used terminology only, while ignoring conceptual patterns. List of weaknesses with using **BOW** (1-3 addressed issues on a lexical level, and 4 conceptual level):

1. Multi-Word Expressions with an own meaning like "European Union" are chunked into pieces with possibly very different meanings like "union".
2. Synonymous Words like "tungsten" and "wolfram" are mapped into different features.
3. Polysemous Words are treated as one single feature while they may actually have multiple distinct meanings.
4. Lack of Generalization: there is no way to generalize similar terms like "beef" and "pork" to their common hypernym "meat".

WordNet database organizes simple words and multi-word expressions of different syntactic categories into so called synonym sets (synsets), each of which represents an underlying concept and links these through semantic relations.

Conceptual Document Representation:

- Candidate Term Detection: Strategy built on the assumption that if you find the longest multi-word expressions in the text, the lexicon will lead to a mapping to the most specific concept for that word (instead of querying single words, which may lead to wrong mapping).
- Syntactical Patterns: Analysis by using POS-tagging.
- Morphological Transformations: Entry form, base form reduction. Stemming if the first query for the inflected forms on the original lexicon turned out unsuccessful.
- Word Sense Disambiguation (WSD): A lexical entry for an expression does not necessarily imply a one-to-one mapping to a concept in the ontology.
- Disambiguate an expression versus multiple possible concepts.
- Generalization: Going from specific concepts in the text to general concept representations. Mapping words based on generalization (up to a certain level).

[5]

QC: predict the entity type of the answer of a natural language question, mostly achieved by using machine learning. Used Latent Semantic Analysis (LSA)

technique to reduce the large feature space of questions to a much smaller and efficient feature space. Two different classifiers: Back-Propagation Neural Networks (BPNN) and Support Vector Machines (SVM). Found that using LSA on question classification made it more time efficient and improved classification accuracy by removing redundant features. Discovered that when the original feature space is compact and efficient, its reduced space performs better than a large feature space with a rich set of features. They also found that in the reduced feature space, BPNN was better than SVM. Competitive with state-of-the-art, even though they used smaller feature space. [23]

Sentiment and such: [27]

Note to self: Map graph over feature impact (unprocessed, singular, all) Also add in estimated training time for exhaustive search (e.g. 120 minutes for SVC vs 100 for SGD over 16k questions (since 4k = test)).

In the paper by Toba et al. [51], they experiment with the use of statistical learning to find the expected answer pattern for factoid QA pairs. E.g. if you ask someone where a certain event took place, the answer pattern would be a location. They group question analysis into two approaches; pattern-based (high precision, low recall) and ML (high recall, low precision¹). Pattern-based would match word sequences against a set of patterns (e.g. regular expressions), whereas ML would be based on the accuracy of the classifier (e.g. lexical or linguistic feature sets). The retrieval of QA pairs is done by using a statistical relation framework: Bayesian Analogical Reasoning (BAR). Features sets are extracted from the training set by use of binary values checking if the question contains a given question word. The BAR framework then learns the related features and computes the estimation for them. Thereafter QA pairs are retrieved from the testing set and compared against the training set. Afterwards, the QA pairs that have identical question words are identified, and overlapping pairs are grouped according their named entity group. To retrieve named entities, they used two different recognizers. The first was Stanford (extracts the person, organization and location), and the second was dictionary based (extract number entities and fine-grained noun-based entities). Question words were extracted by building a question word list from the training set (achieved by using Stanford Part-of-Speech (POS) tagger). Then for each question, look for the appearance of the question word to create the feature set. Mapped named-entities. [51]

Potentially move "all this failed and went wrong" here

To write:

Tutorials that I went through

Using SGD (based on tutorials from scikit-learn)

¹ Low precision can occur if the feature sets are not fitted well enough during classifier training [51, p. 283].

Testing out different text classification algorithms (SVC, SGD and LinearSVC)
Paths, Windows vs. Linux, parallelization, gpu_count, etc

4.1 Data and Testing

discussion on the data set and how it was tested.
the results and what they showed.
potential improvements, etc.

4.2 Artificial Intelligence (AI) Methods

alternative methods and options (e.g. one could have used ann or k-nn, but as
shown in. . .)
not sure if this section is relevant?

4.3 Implementation Architecture

discussion on the code that was written and its functionality
what worked, what should be updated/changed, etc.

4.4 Limitations and other issues

why didn't something work as intended?
why wasn't X completed/implemented?
etc.

5 Conclusion/Summary

5.1 Overview of main results

basically what the title says; a summary of the results

5.2 Further work

additions and updates to the system

new research possibilities based on results

Bibliography

- [1] Saif Ahmed, Seungwon Yang, and Aditya Johri. “Does Online Q&A Activity Vary Based on Topic: A Comparison of Technical and Non-technical Stack Exchange Forums”. In: *Proceedings of the Second (2015) ACM Conference on Learning @ Scale*. L@S ’15. Vancouver, BC, Canada: ACM, 2015, pp. 393–398. ISBN: 978-1-4503-3411-2. DOI: [10.1145/2724660.2728701](https://doi.org/10.1145/2724660.2728701). URL: <http://doi.acm.org/10.1145/2724660.2728701>.
- [2] Ashton Anderson et al. “Discovering Value from Community Activity on Focused Question Answering Sites: A Case Study of Stack Overflow”. In: *Proceedings of the 18th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. KDD ’12. Beijing, China: ACM, 2012, pp. 850–858. ISBN: 978-1-4503-1462-6. DOI: [10.1145/2339530.2339665](https://doi.org/10.1145/2339530.2339665). URL: <http://doi.acm.org/10.1145/2339530.2339665>.
- [3] Jeff Atwood. *A Day in the Penalty Box*. 2009. URL: <http://blog.stackoverflow.com/2009/04/a-day-in-the-penalty-box/> (visited on 04/05/2016).
- [4] Jeff Atwood. *What is Trolling?* 2015. URL: <https://blog.codinghorror.com/what-is-trolling/> (visited on 05/23/2016).
- [5] Stephan Bloehdorn and Andreas Hotho. “Boosting for text classification with semantic features”. In: *WebKDD*. Springer. 2004, pp. 149–166.
- [6] William B Cavnar, John M Trenkle, et al. “N-gram-based text categorization”. In: *Ann Arbor MI* 48113.2 (1994), pp. 161–175.
- [7] CommunityWiki. *Should 'Hi', 'thanks', taglines, and salutations be removed from posts?* 2016. URL: <http://meta.stackexchange.com/questions/2950/should-hi-thanks-taglines-and-salutations-be-removed-from-posts> (visited on 05/07/2016).
- [8] CommunityWiki. *What is a "closed" or "on hold" question?* 2016. URL: <http://meta.stackexchange.com/questions/10582/what-is-a-closed-or-on-hold-question> (visited on 04/05/2016).
- [9] Sebastian Deterding et al. “From game design elements to gamefulness: defining gamification”. In: *Proceedings of the 15th international academic MindTrek conference: Envisioning future media environments*. ACM. 2011, pp. 9–15.

- [10] Eloff. *Strip HTML from strings in Python*. 2009. URL: <http://stackoverflow.com/a/925630> (visited on 03/22/2016).
- [11] Howard Fosdick. *Why People Troll and How to Stop Them*. 2012. URL: <http://www.osnews.com/story/25540> (visited on 05/23/2016).
- [12] Pascal-Emmanuel Gobry. *Google Is Indexing Stack Overflow At 10 Times Per Second*. 2011. URL: <http://www.businessinsider.com/google-stackoverflow-2011-3?r=US&IR=T&IR=T> (visited on 05/23/2016).
- [13] Benjamin V. Hanrahan, Gregorio Convertino, and Les Nelson. “Modeling Problem Difficulty and Expertise in Stackoverflow”. In: *Proceedings of the ACM 2012 Conference on Computer Supported Cooperative Work Companion*. CSCW ’12. Seattle, Washington, USA: ACM, 2012, pp. 91–94. ISBN: 978-1-4503-1051-2. DOI: [10.1145/2141512.2141550](https://doi.org/10.1145/2141512.2141550). URL: <http://doi.acm.org/10.1145/2141512.2141550>.
- [14] Matthias H Heie, Edward WD Whittaker, and Sadaoki Furui. “Question answering using statistical language modelling”. In: *Computer Speech & Language* 26.3 (2012), pp. 193–209.
- [15] Zhiheng Huang, Marcus Thint, and Zengchang Qin. “Question Classification Using Head Words and Their Hypernyms”. In: *Proceedings of the Conference on Empirical Methods in Natural Language Processing*. EMNLP ’08. Honolulu, Hawaii: Association for Computational Linguistics, 2008, pp. 927–936. URL: <http://dl.acm.org/citation.cfm?id=1613715.1613835>.
- [16] Hideki Isozaki. “An Analysis of a High-performance Japanese Question Answering System”. In: 4.3 (Sept. 2005), pp. 263–279. ISSN: 1530-0226. DOI: [10.1145/1111667.1111670](https://doi.org/10.1145/1111667.1111670). URL: <http://doi.acm.org/10.1145/1111667.1111670>.
- [17] Thorsten Joachims. “Machine Learning: ECML-98: 10th European Conference on Machine Learning Chemnitz, Germany, April 21–23, 1998 Proceedings”. In: ed. by Claire Nédellec and Céline Rouveirol. Berlin, Heidelberg: Springer Berlin Heidelberg, 1998. Chap. Text categorization with Support Vector Machines: Learning with many relevant features, pp. 137–142. ISBN: 978-3-540-69781-7. DOI: [10.1007/BFb0026683](https://doi.org/10.1007/BFb0026683). URL: <http://dx.doi.org/10.1007/BFb0026683>.
- [18] Celso Antonio Alves Kaestner. “Support Vector Machines and Kernel Functions for Text Processing”. In: *Revista de Informática Teórica e Aplicada* 20.3 (2013), pp. 130–154.
- [19] Gary Klein. *Blinded By Data*. 2016. URL: <https://www.edge.org/response-detail/26692> (visited on 04/24/2016).

- [20] Stanley Letovsky. “Cognitive processes in program comprehension”. In: *Journal of Systems and Software* 7.4 (1987), pp. 325–339. ISSN: 0164-1212. DOI: [http://dx.doi.org/10.1016/0164-1212\(87\)90032-X](http://dx.doi.org/10.1016/0164-1212(87)90032-X). URL: <http://www.sciencedirect.com/science/article/pii/016412128790032X>.
- [21] C Galina E. Lezina and Artem M. Kuznetsov. *Predict Closed Questions on StackOverflow*. 2013. URL: <http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.394.5678>.
- [22] Xin Li and Dan Roth. “Learning Question Classifiers: The Role of Semantic Information”. In: *Natural Language Engineering* 1.1 (), pp. 000–000.
- [23] B. Loni, S. H. Khoshnevis, and P. Wiggers. “Latent semantic analysis for question classification with neural networks”. In: *Automatic Speech Recognition and Understanding (ASRU), 2011 IEEE Workshop on*. Dec. 2011, pp. 437–442. DOI: [10.1109/ASRU.2011.6163971](https://doi.org/10.1109/ASRU.2011.6163971).
- [24] Vanessa Lopez et al. “Is Question Answering Fit for the Semantic Web?: A Survey”. In: *Semant. web* 2.2 (Apr. 2011), pp. 125–155. ISSN: 1570-0844. DOI: [10.3233/SW-2011-0041](https://doi.org/10.3233/SW-2011-0041). URL: <http://dx.doi.org/10.3233/SW-2011-0041>.
- [25] et al M. Sewak. *Finding a Growth Business Model at Stack Overflow, Inc*. 2010. URL: <https://web.stanford.edu/class/ee204/Publications/Finding%20a%20Growth%20Business%20Model%20at%20Stack%20overflow.pdf> (visited on 05/07/2016).
- [26] Jitendra Maan. “Social business transformation through gamification”. In: *arXiv preprint arXiv:1309.7063* (2013).
- [27] Andrew L. Maas et al. “Learning Word Vectors for Sentiment Analysis”. In: *Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics: Human Language Technologies - Volume 1*. HLT ’11. Portland, Oregon: Association for Computational Linguistics, 2011, pp. 142–150. ISBN: 978-1-932432-87-9. URL: <http://dl.acm.org/citation.cfm?id=2002472.2002491>.
- [28] Christopher D Manning, Prabhakar Raghavan, Hinrich Schütze, et al. *Introduction to information retrieval*. Vol. 1. Cambridge University Press, 2008. ISBN: 978-0521865715.
- [29] Dana Movshovitz-Attias et al. “Analysis of the reputation system and user contributions on a question answering website: Stackoverflow”. In: *Advances in Social Networks Analysis and Mining (ASONAM), 2013 IEEE/ACM International Conference on*. IEEE. 2013, pp. 886–893.

- [30] S. M. Nasehi et al. “What makes a good code example?: A study of programming Q and A in StackOverflow”. In: *Software Maintenance (ICSM), 2012 28th IEEE International Conference on*. Sept. 2012, pp. 25–34. DOI: [10.1109/ICSM.2012.6405249](https://doi.org/10.1109/ICSM.2012.6405249).
- [31] Rodney D Nielsen et al. “A taxonomy of questions for question generation”. In: *Proceedings of the Workshop on the Question Generation Shared Task and Evaluation Challenge*. 2008.
- [32] D. Posnett et al. “Mining Stack Exchange: Expertise Is Evident from Initial Contributions”. In: *Social Informatics (SocialInformatics), 2012 International Conference on*. Dec. 2012, pp. 199–204. DOI: [10.1109/SocialInformatics.2012.67](https://doi.org/10.1109/SocialInformatics.2012.67).
- [33] Stuart J. Russell and Peter Norvig. *Artificial Intelligence: A Modern Approach*. 3rd ed. Pearson Education, 2013. ISBN: 978-1292024202.
- [34] Yutaka Sasaki. “Question answering as question-biased term extraction: a new approach toward multilingual QA”. In: *Proceedings of the 43rd Annual Meeting on Association for Computational Linguistics*. Association for Computational Linguistics. 2005, pp. 215–222.
- [35] Scikitlearn.org. 4.2. *Feature extraction*. 2016. URL: http://scikit-learn.org/stable/modules/feature_extraction.html (visited on 05/07/2016).
- [36] Scikitlearn.org. *sklearn.feature_extraction.text.CountVectorizer*. 2016. URL: http://scikit-learn.org/stable/modules/generated/sklearn.feature_extraction.text.CountVectorizer.html (visited on 05/07/2016).
- [37] Logan Short, Christopher Wong, and David Zeng. “Tag recommendations in stackoverflow”. In: (2014).
- [38] Louisa M. Slowiaczek et al. “Information selection and use in hypothesis testing: What is a good question, and what is a good answer?” In: *Memory & Cognition* 20.4 (1992), pp. 392–405. ISSN: 1532-5946. DOI: [10.3758/BF03210923](https://doi.org/10.3758/BF03210923). URL: <http://dx.doi.org/10.3758/BF03210923>.
- [39] SpaceMachine.net. *DATASETS OVER ALGORITHMS*. 2016. URL: <http://www.spacemachine.net/views/2016/3/datasets-over-algorithms> (visited on 04/24/2016).
- [40] Joel Spolsky. *Stack Overflow Launches*. 2008. URL: <http://www.joelonsoftware.com/items/2008/09/15.html> (visited on 05/06/2016).
- [41] Inc. StackExchange. *Stack Exchange Data Dump*. 2016. URL: <https://archive.org/details/stackexchange> (visited on 04/05/2016).
- [42] StackOverflow.com. *Badges*. 2016. URL: <http://stackoverflow.com/help/badges> (visited on 04/05/2016).

- [43] StackOverflow.com. *How to Ask*. 2016. URL: <http://stackoverflow.com/questions/ask/advice?> (visited on 04/05/2016).
- [44] StackOverflow.com. *Privileges*. 2016. URL: <http://stackoverflow.com/help/privileges> (visited on 04/05/2016).
- [45] StackOverflow.com. *What does it mean if a question is "closed" or "on hold"?* 2016. URL: <http://stackoverflow.com/help/closed-questions> (visited on 04/05/2016).
- [46] StackOverflow.com. *What is reputation? How do I earn (and lose) it?* 2016. URL: <http://stackoverflow.com/help/whats-reputation> (visited on 04/05/2016).
- [47] StackOverflow.com. *What topics can I ask about here?* 2016. URL: <http://stackoverflow.com/help/on-topic> (visited on 04/05/2016).
- [48] StackOverflow.com. *What types of questions should I avoid asking?* 2016. URL: <http://stackoverflow.com/help/dont-ask> (visited on 04/05/2016).
- [49] Clayton Stanley and Michael D Byrne. "Predicting tags for stackoverflow posts". In: *Proceedings of ICCM*. Vol. 2013. 2013.
- [50] Ewan Klein Steven Bird and Edward Loper. *Categorizing and Tagging Words*. 2015. URL: <http://www.nltk.org/book/ch05.html> (visited on 05/06/2016).
- [51] H. Toba, M. Adriani, and R. Manurung. "Expected answer type construction using analogical reasoning in a question answering task". In: *Advanced Computer Science and Information System (ICACSIS), 2011 International Conference on*. Dec. 2011, pp. 283–290.
- [52] C. Treude, O. Barzilay, and M. Storey. "How do programmers ask and answer questions on the web? (NIER track)". In: *Software Engineering (ICSE), 2011 33rd International Conference on*. May 2011, pp. 804–807. DOI: 10.1145/1985793.1985907.
- [53] Bogdan Vasilescu. *Academic papers using Stack Exchange data*. 2012. URL: <http://meta.stackexchange.com/questions/134495/academic-papers-using-stack-exchange-data>.
- [54] Shaowei Wang, David Lo, and Lingxiao Jiang. "An Empirical Study on Developer Interactions in StackOverflow". In: *Proceedings of the 28th Annual ACM Symposium on Applied Computing*. SAC '13. Coimbra, Portugal: ACM, 2013, pp. 1019–1024. ISBN: 978-1-4503-1656-9. DOI: 10.1145/2480362.2480557. URL: <http://doi.acm.org/10.1145/2480362.2480557>.
- [55] Alexander Wissner-Gross. *Datasets Over Algorithms*. 2016. URL: <https://www.edge.org/response-detail/26587> (visited on 04/24/2016).

- [56] Jinzhong Xu, Yanan Zhou, and Yuan Wang. “A Classification of Questions Using SVM and Semantic Similarity Analysis”. In: *Internet Computing for Science and Engineering (ICICSE), 2012 Sixth International Conference on*. Apr. 2012, pp. 31–34. DOI: [10.1109/ICICSE.2012.49](https://doi.org/10.1109/ICICSE.2012.49).
- [57] Jie Yang et al. “User Modeling, Adaptation, and Personalization: 22nd International Conference, UMAP 2014, Aalborg, Denmark, July 7-11, 2014. Proceedings”. In: ed. by Vania Dimitrova et al. Cham: Springer International Publishing, 2014. Chap. Sparrows and Owls: Characterisation of Expert Behaviour in StackOverflow, pp. 266–277. ISBN: 978-3-319-08786-3. DOI: [10.1007/978-3-319-08786-3_23](https://doi.org/10.1007/978-3-319-08786-3_23). URL: http://dx.doi.org/10.1007/978-3-319-08786-3_23.
- [58] Show-Jane Yen et al. “A support vector machine-based context-ranking model for question answering”. In: *Information Sciences* 224 (2013), pp. 77–87. ISSN: 0020-0255. DOI: <http://dx.doi.org/10.1016/j.ins.2012.10.014>. URL: <http://www.sciencedirect.com/science/article/pii/S0020025512006792>.
- [59] Dell Zhang and Wee Sun Lee. “Question Classification Using Support Vector Machines”. In: *Proceedings of the 26th Annual International ACM SIGIR Conference on Research and Development in Informaion Retrieval*. SIGIR '03. Toronto, Canada: ACM, 2003, pp. 26–32. ISBN: 1-58113-646-3. DOI: [10.1145/860435.860443](https://doi.org/10.1145/860435.860443). URL: <http://doi.acm.org/10.1145/860435.860443>.

A Appendix

A.1 Acronyms

AI Artificial Intelligence. [3](#)

BOW Bag of Words. [9](#), [12](#), [15](#), [17](#), [30](#)

IDE Integrated Development Environment. [25](#)

IE Information Extraction. [11](#)

IR Information Retrieval. [10](#), [11](#), [17](#), [19](#)

ML Machine Learning. [iii](#), [2](#), [3](#), [11](#), [20](#), [31](#)

NLTK Natural Language Toolkit. [9](#), [25](#)

POS Part of Speech. [9](#)

QA Question-Answering. [iii](#), [1](#), [6](#), [10–13](#), [31](#)

QC Question Classification. [2](#), [8](#), [12–14](#), [20](#), [30](#)

SE Stack Exchange. [iii](#), [1–4](#), [7](#), [21](#)

SO Stack Overflow. [iii](#), [1–7](#), [28](#)

SVM Support Vector Machines. [ii](#), [iii](#), [2](#), [3](#), [10](#), [11](#), [15](#), [17](#), [19](#), [20](#)

A.2 Data sets/Statistical Overview

A.3 MySQL Database

<div> <div>Posts</div> <div> <div>Id INT</div> <div>PostTypeId INT</div> <div>ParentId INT</div> <div>AcceptedAnswerId INT</div> <div>CreationDate DATETIME</div> <div>Score INT</div> <div>ViewCount INT</div> <div>Body LONGTEXT</div> <div>OwnerUserId INT</div> <div>LastEditorUserId INT</div> <div>LastEditorDisplayName VARCHAR(255)</div> <div>LastEditDate DATETIME</div> <div>LastActivityDate DATETIME</div> <div>CommunityOwnedDate DATETIME</div> <div>ClosedDate DATETIME</div> <div>Title VARCHAR(255)</div> <div>Tags VARCHAR(255)</div> <div>AnswerCount INT</div> <div>CommentCount INT</div> <div>FavoriteCount INT</div> </div> <div>Indexes</div> </div>	<div> <div>negroate_Posts</div> <div> <div>Id INT</div> <div>PostTypeId INT</div> <div>ParentId INT</div> <div>AcceptedAnswerId INT</div> <div>CreationDate DATETIME</div> <div>Score INT</div> <div>ViewCount INT</div> <div>Body LONGTEXT</div> <div>OwnerUserId INT</div> <div>LastEditorUserId INT</div> <div>LastEditorDisplayName VARCHAR(255)</div> <div>LastEditDate DATETIME</div> <div>LastActivityDate DATETIME</div> <div>CommunityOwnedDate DATETIME</div> <div>ClosedDate DATETIME</div> <div>Title VARCHAR(255)</div> <div>Tags VARCHAR(255)</div> <div>AnswerCount INT</div> <div>CommentCount INT</div> <div>FavoriteCount INT</div> </div> <div>Indexes</div> </div>	<div> <div>postvote_Posts</div> <div> <div>Id INT</div> <div>PostTypeId INT</div> <div>ParentId INT</div> <div>AcceptedAnswerId INT</div> <div>CreationDate DATETIME</div> <div>Score INT</div> <div>ViewCount INT</div> <div>Body LONGTEXT</div> <div>OwnerUserId INT</div> <div>LastEditorUserId INT</div> <div>LastEditorDisplayName VARCHAR(255)</div> <div>LastEditDate DATETIME</div> <div>LastActivityDate DATETIME</div> <div>CommunityOwnedDate DATETIME</div> <div>ClosedDate DATETIME</div> <div>Title VARCHAR(255)</div> <div>Tags VARCHAR(255)</div> <div>AnswerCount INT</div> <div>CommentCount INT</div> <div>FavoriteCount INT</div> </div> <div>Indexes</div> </div>	<div> <div>Users</div> <div> <div>Id INT</div> <div>Reputation INT</div> <div>CreationDate DATETIME</div> <div>DisplayName VARCHAR(255)</div> <div>EmailHash VARCHAR(255)</div> <div>LastAccessDate DATETIME</div> <div>WebsiteUrl VARCHAR(45)</div> <div>Location VARCHAR(255)</div> <div>Age INT</div> <div>AboutMe VARCHAR(255)</div> <div>Views INT</div> <div>Upvotes INT</div> <div>Downvotes INT</div> </div> <div>Indexes</div> </div>		
<div> <div>PostHistory</div> <div> <div>Id INT</div> <div>PostHistoryTypeId INT</div> <div>PostId VARCHAR(45)</div> <div>RevisionGUID VARCHAR(255)</div> <div>CreationDate DATETIME</div> <div>UserId INT</div> <div>UserIdInNew VARCHAR(255)</div> <div>Text LONGTEXT</div> <div>CloseReasonId INT</div> </div> <div>Indexes</div> </div>	<div> <div>Comments</div> <div> <div>Id INT</div> <div>PostId INT</div> <div>Score INT</div> <div>Text LONGTEXT</div> <div>CreationDate DATETIME</div> <div>UserId INT</div> </div> <div>Indexes</div> </div>	<div> <div>Tags</div> <div> <div>Id INT</div> <div>TagName VARCHAR(255)</div> <div>ExceptPostId INT</div> <div>WikiPostId INT</div> </div> <div>Indexes</div> </div>	<div> <div>Votes</div> <div> <div>Id INT</div> <div>PostId INT</div> <div>VoteTypeId INT</div> <div>CreationDate DATETIME</div> <div>UserId INT</div> <div>BountyAmount INT</div> </div> <div>Indexes</div> </div>	<div> <div>PostLinks</div> <div> <div>Id INT</div> <div>CreationDate DATETIME</div> <div>PostId INT</div> <div>RelatedPostId INT</div> <div>PostLinkTypeId INT</div> </div> <div>Indexes</div> </div>	<div> <div>Badges</div> <div> <div>UserId INT</div> <div>Name VARCHAR(255)</div> <div>Date DATETIME</div> </div> <div>Indexes</div> </div>

Figure 4: MySQL Database used for dataset